


AIRDROP TRAJECTORY OPTIMISATION FOR DISASTER RELIEF

Guiding Supply Packages to a Target in Unknown Wind

ATOMIC ENERGY CENTRAL SCHOOL, MYSURU



BY: AMULYA D

Interactive version: <https://1101-hub.github.io/guided-airdrop/>

Abstract

Relief supplies airdropped to flood-isolated communities drift with the wind and land away from the intended point. This work simulates a self-steering parafoil carrying no wind sensor and asks what bounds its accuracy. Dubins path planning is combined with a law producing a path of prescribed rather than minimum length; an Extended Kalman Filter (EKF) estimates wind, airspeed and heading bias from satellite velocity. Over $N = 150$ randomised drops the median miss distance falls from 8.57 m without a wind estimate to 3.89 m with the EKF (95% CI 2.85–5.28), a paired median reduction of 4.26 m. Turn radius scales linearly with wing loading and the loading needed to exceed wind speed W scales as W^2 , so the error floor scales as W^2 .

1 Introduction

A cloudburst on 19 July 2026 flooded four districts of Upper Assam; by 25 July, 705,000 people were affected and 80,000 remained unreachable (Tribune India, 2026; Onmanorama, 2026). Indian Air Force helicopters airdropped supplies, and the Chief Minister stated an intention to attempt “targeted drops in larger quantities” (ANI, 2026). A payload released without guidance is advected by the air mass throughout its descent, so a requirement for targeted delivery implies current delivery is untargeted.

A parafoil is a ram-air canopy that glides, so horizontal displacement can be commanded during descent. Guided parafoils are established technology: the Naval Postgraduate School *Snowflake* has landed within 3 m of target (Naval Postgraduate School, n.d.), and the open-source ParaDrone plans Dubins paths on microcontroller hardware (Daniel, n.d.). This work claims no priority over either; it supplies a design analysis they do not report, relating wing loading and wind speed to the bound on landing error.

2 Methods

Flight model. As a point mass in steady glide with lift equal to weight, the parafoil has airspeed

$$v = \sqrt{\frac{2g}{\rho C_L}} \sqrt{\frac{m}{S}}, \quad (1)$$

with ρ air density, C_L lift coefficient, m payload mass, S canopy area; only the wing loading m/S appears. When wind exceeds v the vehicle is advected downwind whatever the control input. Values used were $C_L = 0.80$, $C_D = 0.27$ and 20° bank.

Path planning. Under a minimum turn radius the shortest route between oriented configurations is a Dubins path: at most three turning or straight segments in one of six orderings (Dubins, 1957). All six closed forms were verified against step-wise integration (largest endpoint disagreement 9.0×10^{-12} m over 3,514 solutions).

```
def solve(start, end, R, only=None):
    """Shortest Dubins path: try all six shapes, keep the shortest."""
    d = np.hypot(x1 - x0, y1 - y0) / R          # normalise into units of R
    phi = np.arctan2(y1 - y0, x1 - x0)
    a, b = _mod2pi(th0 - phi), _mod2pi(th1 - phi)

    best = None
    for name in (only or TYPES):                  # LSL LSR RSL RSR RLR LRL
        out = _SOLVERS[name](a, b, d)
        if out is None:                            # cannot connect these two
            continue
        t, p, q = out
        length = (t + p + q) * R
        if best is None or length < best["length"]:
```

```

        best = {"type": name, "length": length, "R": R}
    return best

```

Listing 1: Path planning, from `model/dubins.py`. An infinite question — which curve is best? — becomes picking the smallest of six numbers. A shape that cannot connect the two configurations returns `None` and is skipped.

Path-length control. Sink rate is constant, so time aloft and distance flown are fixed at release: the path needs prescribed, not minimum, length. A track at angle α to the bearing extends path length by $1/\cos \alpha$, giving the law $\cos \alpha = d/s$ for remaining distance d and budget s ; α alternates in sign. As altitude falls, $d/s \rightarrow 1$ and $\alpha \rightarrow 0$, giving a straight final approach with no mode change. Excess altitude is spent in a holding circle.

```

ratio = d_eff / budget          # Dubins length, not straight-line distance

if ratio < np.cos(self.alpha_max):
    # ORBIT: too high even for the steepest weave, so make no net progress
    radial_err = np.clip((d - R) / max(R, 1e-6), -1.0, 1.0)
    desired = bearing + self.side * (np.pi/2.0 - radial_err * np.pi/2.0)
else:
    # WEAVE: cos(alpha) = d / budget, alternating side to side
    alpha = np.arccos(np.clip(ratio, -1.0, 1.0))
    err = _wrap(th - bearing)
    if self.side > 0 and err > alpha:      self.side = -1.0
    elif self.side < 0 and err < -alpha:   self.side = 1.0
    desired = bearing + self.side * alpha

turn = np.sign(_wrap(desired - th))

```

Listing 2: The control law, from `model/guidance.py`. `d_eff` is the Dubins length rather than the straight-line distance, so the turn needed to line up is charged to the budget. Using the straight line made the vehicle believe it had height to spare, and it landed short.

Wind estimation. No wind sensor is carried. An EKF (Kalman, 1960) maintains wind components w_x , w_y , airspeed v_a and magnetometer bias b from GNSS ground velocity, nonlinearity entering through b in the heading trigonometry. A headwind and an airspeed underestimate are indistinguishable on a fixed heading, leaving the state unobservable; turning separates them, and path-length control already turns.

```

wx, wy, va, b = self.x
th = heading_meas + b                # b is the compass bias: the
c, s = np.cos(th), np.sin(th)       # nonlinearity that forces an EKF

h = np.array([va * c + wx, va * s + wy]) # predicted ground velocity
H = np.array([[1.0, 0.0, c, -va * s],    # Jacobian dh/dx
              [0.0, 1.0, s, va * c]])

y = np.asarray(v_ground_meas, float) - h # innovation
S = H @ self.P @ H.T + self.R
K = self.P @ H.T @ np.linalg.inv(S)     # Kalman gain
self.x = self.x + K @ y


```

Listing 3: The measurement update, from `model/ekf.py`. The bias b sits inside the heading trigonometry, which is what makes the problem nonlinear and the filter an *extended* one; with b known the same problem would be linear.

Simulation. Descent is integrated at $\Delta t = 0.05$ s, replanning every step, and guidance sees only GNSS velocity with Gaussian noise of 0.25 m s^{-1} and a magnetometer bias from $\pm 12^\circ$, never the simulated wind. Each trial draws wind $0\text{--}3.2 \text{ m s}^{-1}$ with uniform heading and release height $20\text{--}60$ m, flown three times under identical conditions: *no-estimate* (still air assumed), *EKF* (wind estimated onboard) and *true-*

wind (simulated wind supplied). *Miss distance* is horizontal separation between touchdown and target; intervals are bootstrap percentiles from 20,000 resamples.

3 Results



(a) No-estimate configuration. Median miss distance 15.03 m.

(b) EKF configuration. Median miss distance 2.26 m.

Figure 1: Ground tracks of $N = 8$ descents, plotted in a target-centred frame with east and north in metres on identical axes in both panels. The same eight sets of conditions are flown in each panel, with wind speed drawn from $1.0\text{--}3.0\text{ m s}^{-1}$ and release height from $30\text{--}55\text{ m}$. The filled circle at the origin marks release, the cross marks the target, and filled circles at track ends mark touchdown. Touchdown points are dispersed in fig. 1a and concentrated near the target in fig. 1b.

Table 1 gives miss distance over $N = 150$ randomised drops. Compared trial by trial, the median reduction from adding the EKF is 4.26 m (95% CI $3.08\text{--}6.57$), and the EKF gives the smaller miss distance in 112 of 150 trials; wind velocity is recovered to a median error of 0.104 m s^{-1} ($0.096\text{--}0.112$). The EKF and true-wind intervals overlap over most of their range, so those configurations are not separated at this sample size.

Table 1: Miss distance by guidance configuration over $N = 150$ randomised drops. Conditions are identical across rows and only the wind information available to the guidance differs. Intervals on the median are 95% bootstrap percentile intervals; the uncertainty on the mean is the standard error.

Configuration	Median (m)	Mean (m)	90th percentile (m)
No-estimate	8.57 (6.49–12.46)	11.52 ± 0.81	24.27
EKF	3.89 (2.85–5.28)	4.94 ± 0.31	10.75
True-wind	4.62 (3.37–5.44)	4.84 ± 0.32	10.93

Two scaling relations follow from eq. (1). Turn radius $R = v^2/(g \tan \phi)$ at bank angle ϕ , with $v \propto \sqrt{m/S}$, gives $R \propto m/S$: across loadings $0.3\text{--}2.4\text{ kg m}^{-2}$ the ratio $R/(m/S)$ is $4.7694\text{ m}^3\text{ kg}^{-1}$, constant to $1.9 \times 10^{-14}\%$. Wind penetration requires $v > W$, giving $(m/S)_{\min} \propto W^2$: for $W = 2.0\text{--}4.0\text{ m s}^{-1}$

the ratio $(m/S)_{\min}/W^2$ is $0.058722 \text{ kg m}^{-2}(\text{m s}^{-1})^{-2}$, constant to $1.2 \times 10^{-14}\%$. Both sit at double-precision level: these are algebraic identities, not empirical fits. Composing them gives

$$R_{\min} \propto W^2. \quad (2)$$

Figure 2 maps miss distance over this space.




Figure 2: Median miss distance over wing loading m/S and wind speed W , from 24 drops at each of 13×11 grid points in the EKF configuration at a release height of 30 m. Contours are labelled in metres. The line marks $v = W$; in the shaded region below and left of it airspeed does not exceed wind speed and no guided descent is possible. Miss distance is smallest at low wing loading and low wind speed, and exceeds 12 m at the highest loadings irrespective of wind speed.

4 Discussion

The paired reduction establishes that onboard wind estimation improves accuracy here. That the EKF is not separated from true-wind is consistent with its correcting two error sources where true-wind corrects one: supplying wind leaves the $\pm 12^\circ$ magnetometer bias uncorrected, which the filter estimates jointly. Separating them would need a larger sample.

Equation (2) follows from two competing requirements: penetrating wind of speed W demands a wing loading rising as W^2 , while turn radius rises linearly with loading and bounds the smallest correction achievable in the final seconds. Specifying for wind never encountered therefore degrades accuracy in all conditions, including calm. The appropriate choice is the lowest loading satisfying $v > W$ with margin; set by ballast, it can be matched to the forecast before release, unlike canopy area.

Three limitations bound these results: turbulence is absent, so the wind field is uniform and steady and the accuracy reported is an upper bound; canopy deformation is excluded, fixing C_L and C_D ; and actuator dynamics are excluded, so turns begin instantaneously. Quantifying this optimism requires flight measurement; the model specifies a test configuration — 1 m² canopy, 0.5–0.8 kg payload, 20 m minimum release height — so it can be falsified rather than fitted.

References

- ANI. (2026, July 25). *IAF choppers deployed to airdrop essential relief supplies in flood-affected areas: Assam CM Himanta Biswa Sarma*.
- Daniel, K. (n.d.). *ParaDrone: Autopilot for parachutes*. <https://hackaday.io/project/176779-paradrone-autopilot-for-parachutes>
- Dubins, L. E. (1957). On curves of minimal length with a constraint on average curvature, and with prescribed initial and terminal positions and tangents. *American Journal of Mathematics*, 79(3), 497–516.
- Kalman, R. E. (1960). A new approach to linear filtering and prediction problems. *Transactions of the ASME—Journal of Basic Engineering*, 82(D), 35–45.
- Naval Postgraduate School. (n.d.). *Snowflake*. Aerodynamic Decelerator Systems Laboratory. <https://nps.edu/web/adsc/snowflake>
- Onmanorama. (2026, July 25). *Assam flood situation remains grim; toll rises to 61, over 7 lakh affected*.
- Tribune India. (2026, July 25). *Govt yet to reach 80,000 flood victims, says Assam CM; toll rises to 61*.